




3fs

# Containers: A Peek Under the Hood

Aleš Brelih

Backend software engineer

13 APRIL 2024



# A little bit about myself



R&D engineer

@ <https://www.arctur.si/>



**Backend software engineer**

@ <https://www.3fs.cloud>

## How did I get here?

- I've been actively working with containers for over 7 years.
- For much of that time, I didn't really understand what was happening under the hood. Tools like Docker and Kubernetes really do a good job at hiding the complexity.
- My interest in security made me dig deeper, leading to a realization: you can't protect what you're unaware of.

# What is a container?

# Let's google it



## Docker

"A container is a standard unit of software that packages up code and all its dependencies so the application runs quickly and reliably from one computing environment to another."



## Google Cloud

"Containers are lightweight packages of software that contain all of the necessary elements to run in any environment."



## Red Hat

"Containers are technologies that allow the packaging and isolation of applications with their entire runtime environment—all of the files necessary to run."

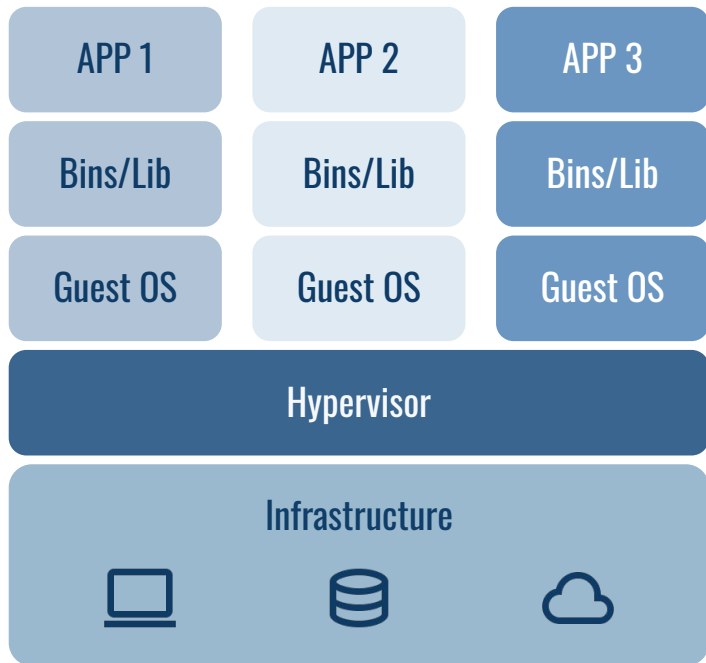
## Probably most heard answer?

“Oh you mean docker? It is something like a Virtual Machine but lightweight.”

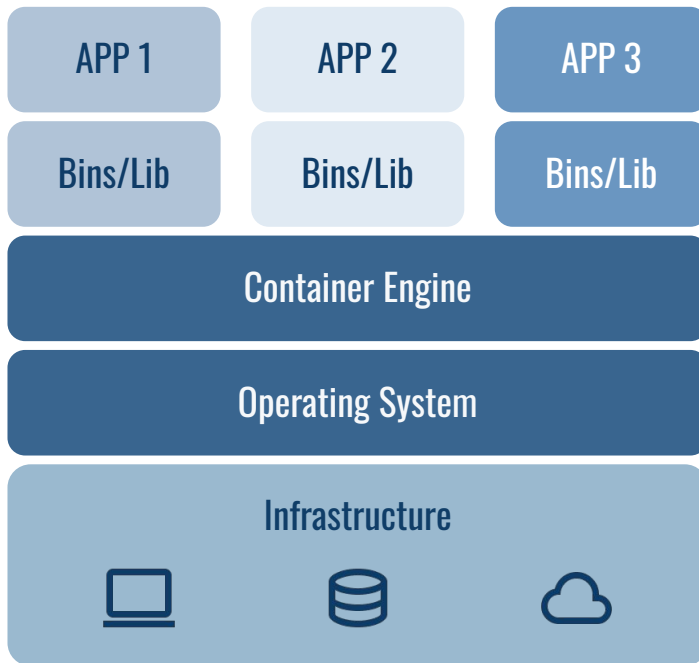
# Truth?

- They are **just processes**, just like any other application that runs directly on the host
- The major difference between containers and VMs is that containers **share host kernel** while VMs virtualize hardware and run multiple guest operating systems on a single physical machine.

## Virtual Machines




## Containers





## Why does the shared kernel matter?



# DEMO

It is just a process.

# What makes a process a container?



## Isolation

Processes should run independently, unaware of other processes or the host system.



## Encapsulation

Wraps everything an application needs to run (including dependencies and libraries) into a single package.



## Resource restriction

Controls and limits the amount of CPU, memory, and disk resources that each container can use.

# Namespaces

Isolation

# Namespaces

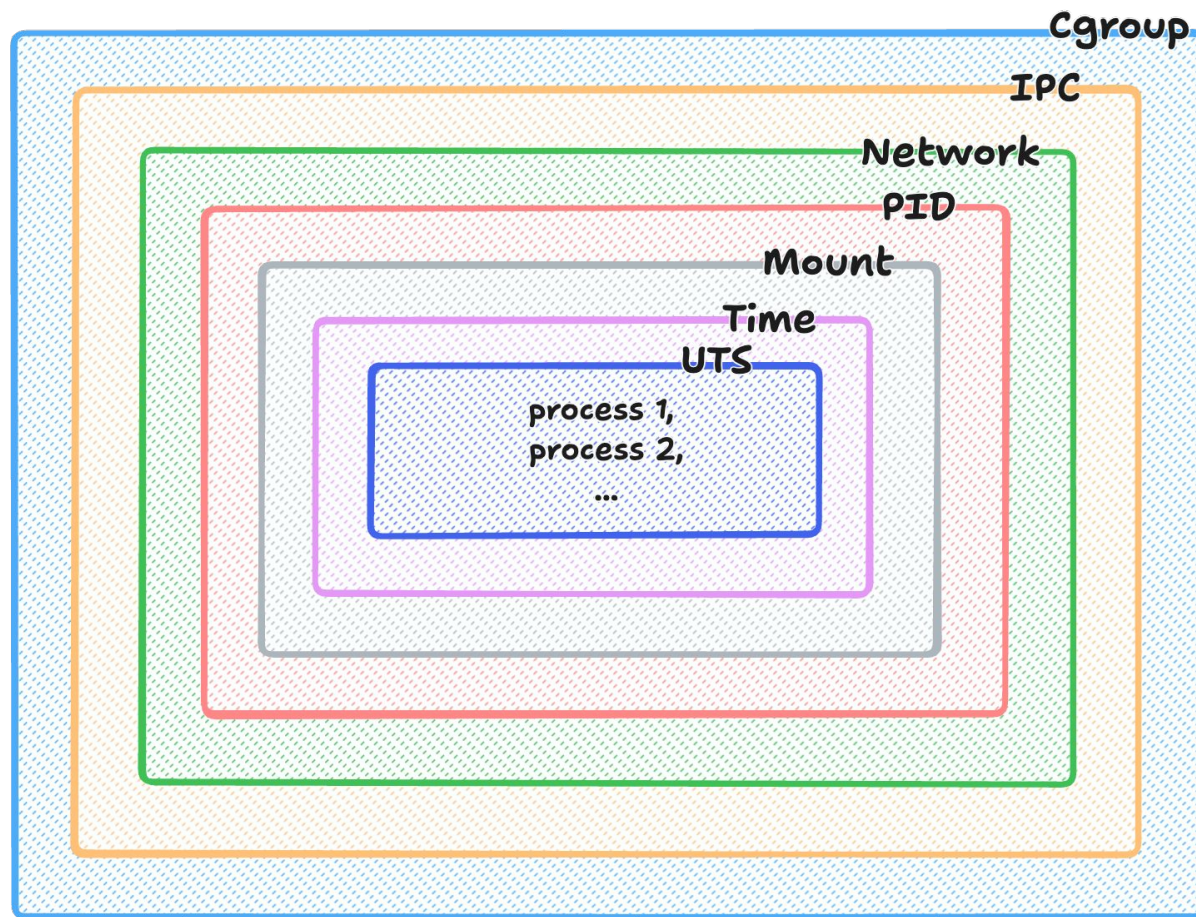
- Allow a process to have its own isolated instance of global resources (e.g., process IDs, network interfaces).
- They limit the potential impact of malicious processes.
- Changes to the system resource are visible to other processes that are members of the namespace, but are invisible to other processes.

# Types

- Control Groups (Cgroups)
- Inter Process Communication (IPC)
- Network
- Mount
- Process ID (PID)
- Time
- User
- Unix Time Sharing (UTS)

3fs

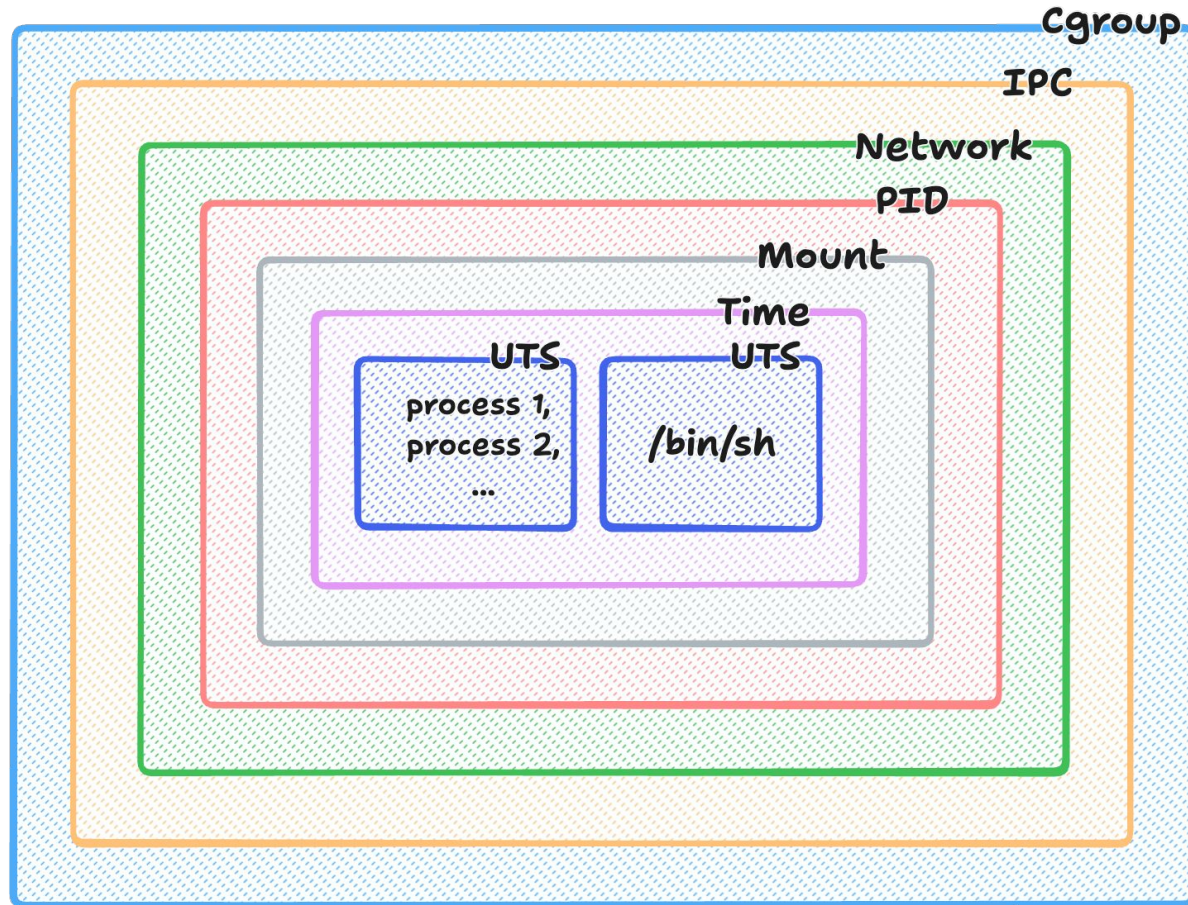
## Default namespace layout





3fs

## What we want to do in demo





# DEMO


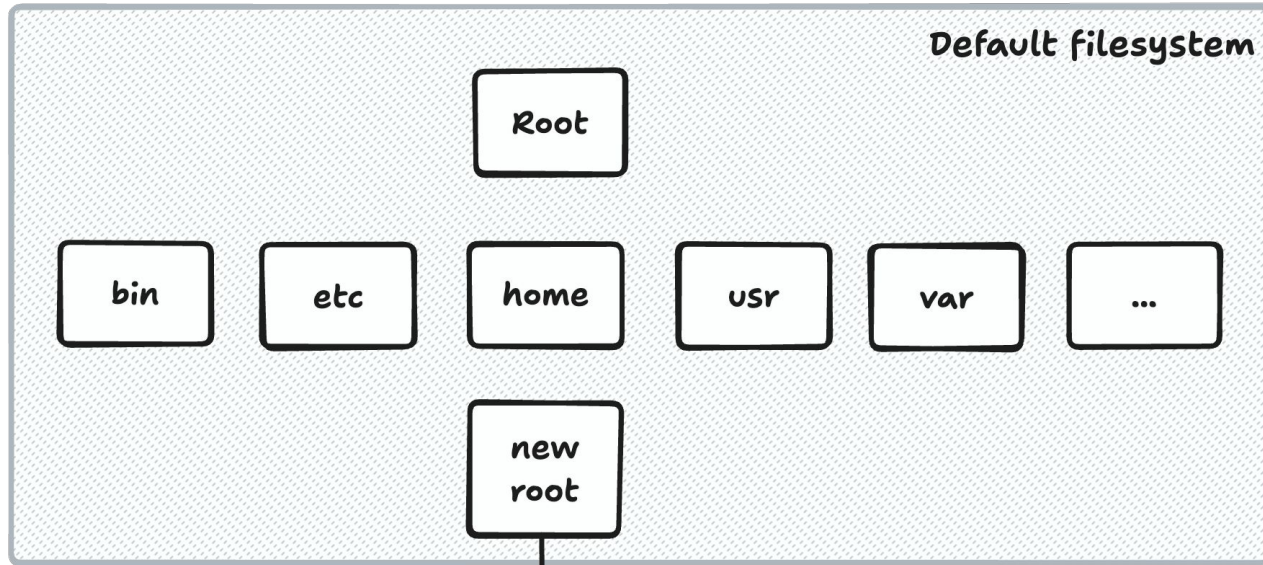
- Unix Time Sharing namespace
- Process ID namespace

# Chroot

Encapsulation

# Chroot

- Used to run command or interactive shell with a new root directory.



# DEMO

- Chroot.

# Docker image layering

```
FROM node:alpine3.19


WORKDIR /app

COPY ./package.json ./package.json

RUN npm install
```

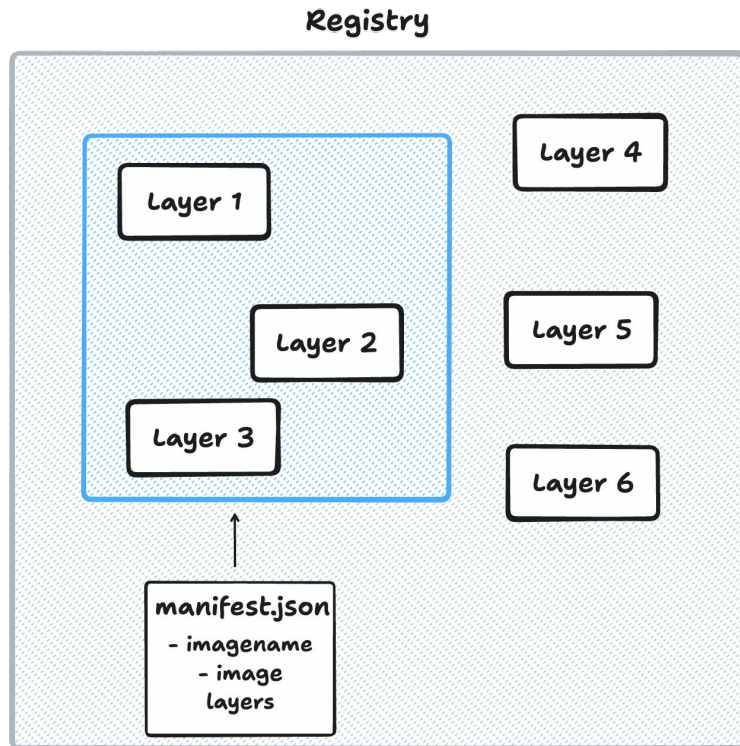


*customimage*



## Layers inside container registry

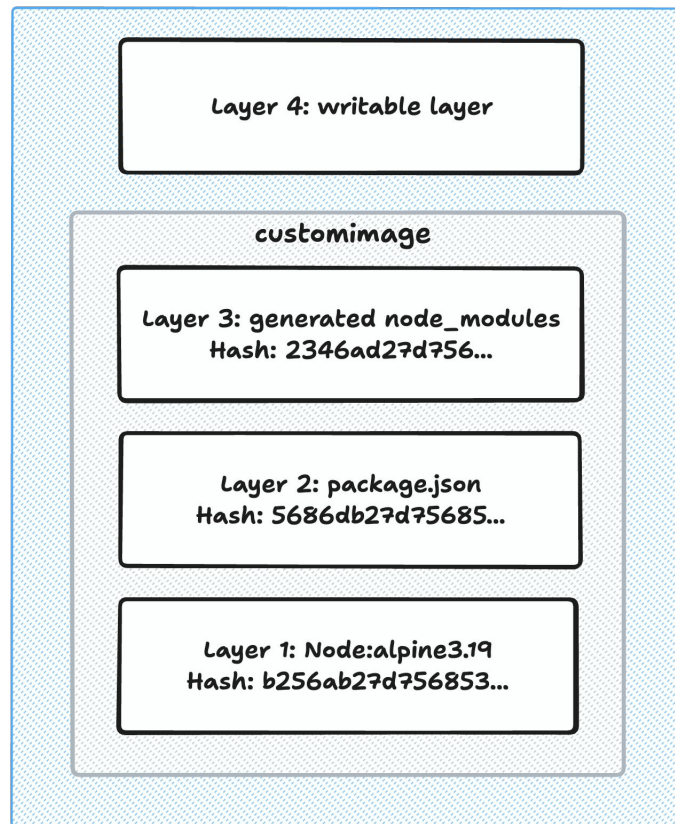
- Layers are independent entities in a Docker registry.
- Manifest.json defines the composition of a Docker image



## Once used

- A new writable layer on top of the image's read-only layers is created.
- Allowing for the original image to remain unchanged and reusable.

### customimage container





# DEMO


- Docker image inspection.

# Controls Groups

Resource limitations

# Cgroups


- Allow you to allocate resources (CPU time, system memory, network bandwidth, or combinations of these resources among processes)
- Cgroups are organized hierarchically (child cgroups inherit some of the attributes of their parents)



# DEMO

- Cgroups in actions.

## Let's put everything together.



# Container escapes

And how to be “safer”

## Be careful when using

- **-privileged** flag (124k files on Github)
- **-v /var/run/docker.sock:/var/run/docker.sock** mount (171k files on Github)
- **-cap-add=\*** flags (52 files on Github)

# DEMO

- Escape from a privileged container.
- Escape out of container with mounted docker.sock.





3fs

**QUESTIONS?**

